

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix, roc_auc_score, roc_curve
```

## get data :

```
ad_data = pd.read_csv('C:\\Users\\DELL\\OneDrive\\Bureau\\
selfeducations\\projects\\Logistic Regression advertising\\
advertising.csv')
```

*# show the head of data 5 first rows*

```
ad_data.head()
```

	Daily Time Spent on Site	Age	Area Income	Daily Internet Usage	\
0	68.95	35	61833.90	256.09	
1	80.23	31	68441.85	193.77	
2	69.47	26	59785.94	236.50	
3	74.15	29	54806.18	245.89	
4	68.37	35	73889.99	225.58	

	Ad Topic Line	City	Male
Country \			
0	Cloned 5thgeneration orchestration	Wrightburgh	0
Tunisia			
1	Monitored national standardization	West Jodi	1
Nauru			
2	Organic bottom-line service-desk	Davidton	0
San Marino			
3	Triple-buffered reciprocal time-frame	West Terrifurt	1
Italy			
4	Robust logistical utilization	South Manuel	0
Iceland			

	Timestamp	Clicked on Ad
0	2016-03-27 00:53:11	0
1	2016-04-04 01:39:02	0
2	2016-03-13 20:35:42	0
3	2016-01-10 02:31:19	0
4	2016-06-03 03:36:18	0

*# show the shape of data :*

```
ad_data.shape
```

```
(1000, 10)
```

```
#show only columns names :
ad_data.columns

Index(['Daily Time Spent on Site', 'Age', 'Area Income',
      'Daily Internet Usage', 'Ad Topic Line', 'City', 'Male',
      'Country',
      'Timestamp', 'Clicked on Ad'],
      dtype='object')
```

```
# show data columns and data type :
ad_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Daily Time Spent on Site              1000 non-null   float64
1   Age                                    1000 non-null   int64
2   Area Income                           1000 non-null   float64
3   Daily Internet Usage                  1000 non-null   float64
4   Ad Topic Line                         1000 non-null   object
5   City                                  1000 non-null   object
6   Male                                  1000 non-null   int64
7   Country                               1000 non-null   object
8   Timestamp                             1000 non-null   object
9   Clicked on Ad                         1000 non-null   int64
dtypes: float64(3), int64(3), object(4)
memory usage: 78.2+ KB
```

```
# descriptive statistic :
ad_data.describe()
```

	Daily Time Spent on Site	Age	Area Income	\
count	1000.000000	1000.000000	1000.000000	
mean	65.000200	36.009000	55000.000080	
std	15.853615	8.785562	13414.634022	
min	32.600000	19.000000	13996.500000	
25%	51.360000	29.000000	47031.802500	
50%	68.215000	35.000000	57012.300000	
75%	78.547500	42.000000	65470.635000	
max	91.430000	61.000000	79484.800000	

	Daily Internet Usage	Male	Clicked on Ad
count	1000.000000	1000.000000	1000.000000
mean	180.000100	0.481000	0.50000
std	43.902339	0.499889	0.50025
min	104.780000	0.000000	0.00000
25%	138.830000	0.000000	0.00000
50%	183.130000	0.000000	0.50000


75%	218.792500	1.000000	1.00000
max	269.960000	1.000000	1.00000

## Exploratory Data Analysis :

### Univariate Analysis:

#### Distribution of age :


```
plt.figure(figsize=(8,5)) # set a fig size 8 by 5
sns.histplot(ad_data['Age'],bins= 20) # creat a histgrame of 20 bins
of Age
plt.title('Distribution of Age')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()
```



#### Distribution of Gender :

```
plt.figure(figsize=(8,5))
sns.countplot(x='Male',data = ad_data) # Create a count plot with
```

```
'Male' as the x-axis and 'ad_data' as the data source
plt.xlabel('Gender')
plt.ylabel('Count')
plt.title('Gender Distribution')
plt.xticks([0, 1], ['Female', 'Male']) # Set custom tick labels for
the x-axis
plt.show()
```




## Correlation :

```
correlation_matrix = ad_data.corr()
plt.figure(figsize = (10,8))
sns.heatmap(correlation_matrix,annot = True, cmap = 'coolwarm',fmt =
'2f' )
plt.title('Correlation Matrix')
plt.show()
```

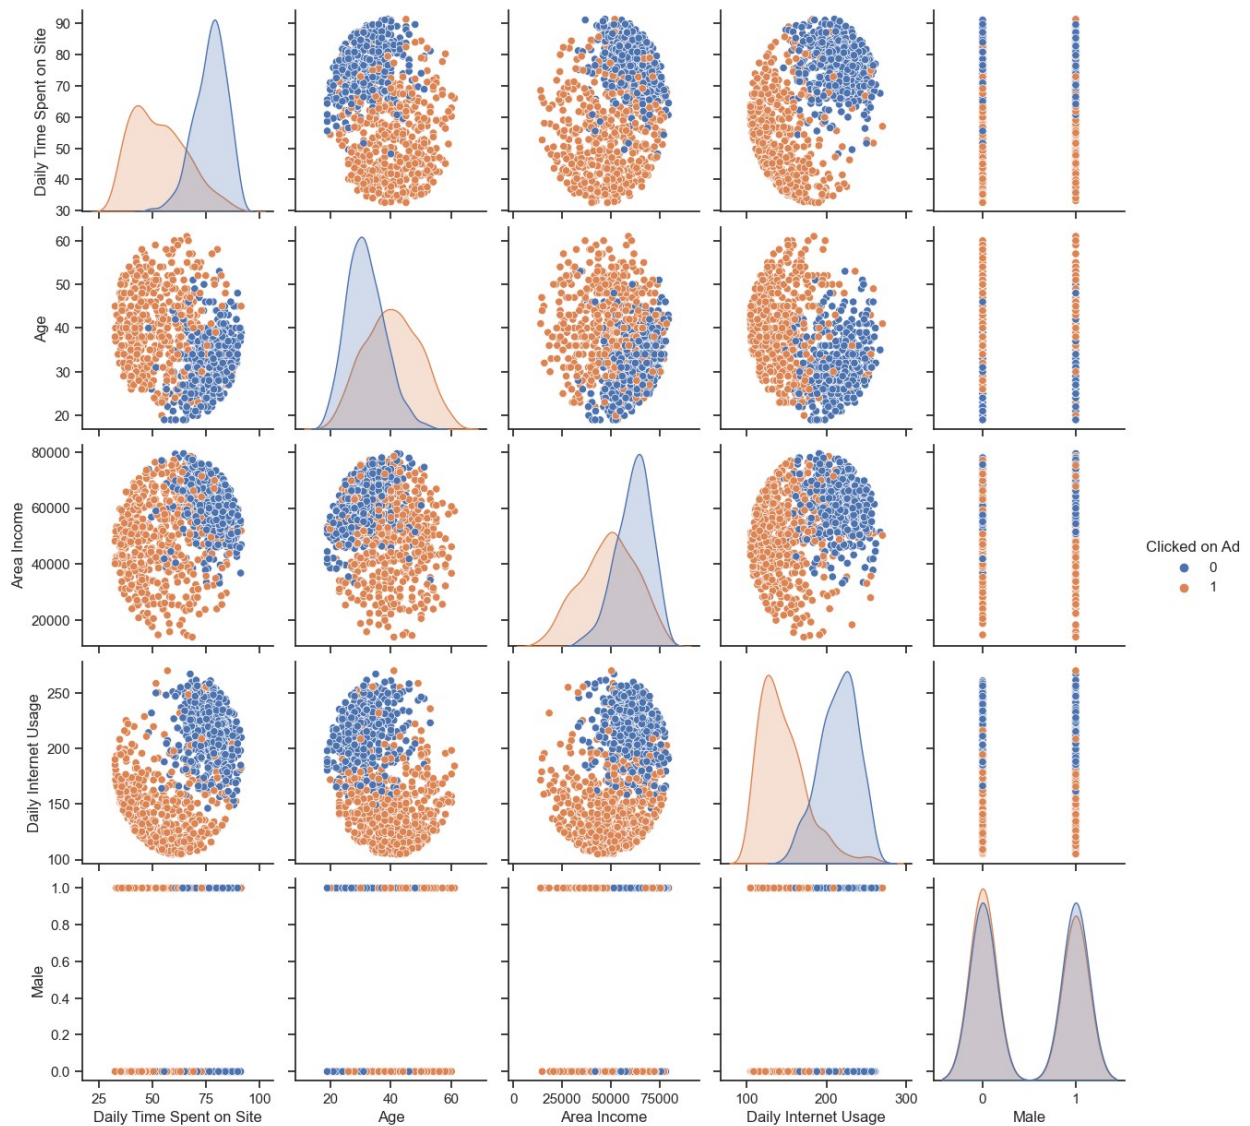
C:\Users\DELL\AppData\Local\Temp\ipykernel\_7420\1951576788.py:1:  
FutureWarning: The default value of numeric\_only in DataFrame.corr is  
deprecated. In a future version, it will default to False. Select only  
valid columns or specify the value of numeric\_only to silence this  
warning.

```
correlation_matrix = ad_data.corr()
```



Vizualization of relations betewwen colmunns:

```
sns.pairplot(ad_data,hue='Clicked on Ad')
<seaborn.axisgrid.PairGrid at 0x1e85c8a58a0>
```



# Logistique Regression :

## Data Preprocessing

```
# create the input variables for data
X = ad_data[['Daily Time Spent on Site', 'Age', 'Area Income', 'Daily
Internet Usage', 'Male']]
# create the target variable click on add as output :
Y = ad_data['Clicked on Ad']

# Split the data into train data 80% and test data 20% :
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=20)
```

## Model Training

```
# initialize the Logistic Regression Model :
logReg_model = LogisticRegression()
# train the model :
logReg_model.fit(X_train,Y_train)

LogisticRegression()
```

## Model Evaluation :

```
# make prediction on the test set :
Y_pred = logReg_model.predict(X_test)

results_df = pd.DataFrame({'Y_pred': Y_pred, 'Y_test': Y_test})
results_df.head()
```

	Y_pred	Y_test
890	0	0
694	0	0
798	1	0
147	1	1
858	1	1

```
# calculate the accuracy of the model :
accuracy = accuracy_score(Y_test,Y_pred)
print("The Accuracy = ",accuracy)
print("The Accuracy = ",accuracy*100,'%')
```

```
The Accuracy = 0.91
The Accuracy = 91.0 %
```

```
# generate a classification report :
Classification_rep =
classification_report(Y_test,Y_pred,target_names=['Classe 0 ','Classe
1'],output_dict=True)
#convert the classification report dictionary to DataFrame:
report_df = pd.DataFrame(Classification_rep).transpose()
print("Classification report : \n\n")
report_df.head()
```

Classification report :


	precision	recall	f1-score	support
Classe 0	0.934579	0.900901	0.917431	111.00
Classe 1	0.881720	0.921348	0.901099	89.00
accuracy	0.910000	0.910000	0.910000	0.91

```
macro avg    0.908150  0.911125  0.909265  200.00
weighted avg 0.911057  0.910000  0.910163  200.00
```

```
# Generate a confusion matrix :
conf_matrix = confusion_matrix(Y_test,Y_pred)
# transform it to DataFrame :
conf_matrix_df = pd.DataFrame(conf_matrix).transpose()
conf_matrix_df
```

```
      0    1
0   100    7
1    11   82
```

```
# Visualize the Confusion Matrix
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix, annot=True, cmap='Blues', fmt='d')
plt.title('Confusion Matrix')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```



```
#calculate the roc AUC score :
roc_auc_score =
roc_auc_score(Y_test,logReg_model.predict_proba(X_test)[:,-1])
print('ROC AUC Score : ', roc_auc_score)
```



ROC AUC Score : 0.9580929243850592

ROC AUC score of 0.958 indicates a highly effective classifier with a strong ability to discriminate between the two classes in the binary classification problem. It's considered a very good result and suggests that the model is making accurate predictions for the target variable.

```
# Additional Visualization (ROC Curve) :  
fpr, tpr, thresholds = roc_curve(Y_test,  
logReg_model.predict_proba(X_test)[: , 1])  
plt.figure(figsize=(8, 6))  
plt.plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc_score:.2f})')  
plt.plot([0, 1], [0, 1], 'k--')  
plt.xlabel('False Positive Rate')  
plt.ylabel('True Positive Rate')  
plt.title('ROC Curve')  
plt.legend()  
plt.show()
```

